

# Activity 1.2 : Training Neural Networks

Barcelona, Jacob Seth S.

Instructor: Engr. Roman Richard

## Objective(s):

This activity aims to demonstrate how to train neural networks using keras

## Intended Learning Outcomes (ILOs):

- Demonstrate how to build and train neural networks
- Demonstrate how to evaluate and plot the model using training and validation loss

## Resources:

- Jupyter Notebook

CI Pima Diabetes Dataset

- pima-indians-diabetes.csv

## Procedures

Load the necessary libraries

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, precision_recall_curve, roc_auc_s
from sklearn.ensemble import RandomForestClassifier

import seaborn as sns

%matplotlib inline
```

```
In [ ]: ## Import Keras objects for Deep Learning
```

```
from keras.models import Sequential
from keras.layers import Input, Dense, Flatten, Dropout, BatchNormalization
from keras.optimizers import Adam, SGD, RMSprop
```

Load the dataset

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [ ]: filepath = "/content/drive/MyDrive/Datasci3/pima-indians-diabetes.csv"
names = ["times_pregnant", "glucose_tolerance_test", "blood_pressure", "skin_thickn",
         "bmi", "pedigree_function", "age", "has_diabetes"]
diabetes_df = pd.read_csv(filepath, names=names)
```

Check the top 5 samples of the data

```
In [ ]: print(diabetes_df.shape)
diabetes_df.sample(5)
```

(768, 9)

```
Out[5]:
```

	times_pregnant	glucose_tolerance_test	blood_pressure	skin_thickness	insulin	bmi	pedigree_function
565	2	95	54	14	88	26.1	
656	2	101	58	35	90	21.8	
504	3	96	78	39	0	37.3	
186	8	181	68	36	495	30.1	
271	2	108	62	32	56	25.2	

```
In [ ]: diabetes_df.dtypes
```

```
Out[6]: times_pregnant      int64
glucose_tolerance_test    int64
blood_pressure            int64
skin_thickness            int64
insulin                  int64
bmi                      float64
pedigree_function         float64
age                      int64
has_diabetes              int64
dtype: object
```

```
In [ ]: X = diabetes_df.iloc[:, :-1].values
        y = diabetes_df["has_diabetes"].values
```

Split the data to Train, and Test (75%, 25%)

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

```
In [ ]: np.mean(y), np.mean(1-y)
```

```
Out[9]: (0.3489583333333333, 0.6510416666666666)
```

Build a single hidden layer neural network using 12 nodes. Use the sequential model with single layer network and input shape to 8.

Normalize the data

```
In [ ]: normalizer = StandardScaler()
        X_train_norm = normalizer.fit_transform(X_train)
        X_test_norm = normalizer.transform(X_test)
```

Define the model:

- Input size is 8-dimensional
- 1 hidden layer, 12 hidden nodes, sigmoid activation
- Final layer with one node and sigmoid activation (standard for binary classification)

```
In [ ]: model = Sequential([
        Dense(12, input_shape=(8,)), activation="relu"),
        Dense(1, activation="sigmoid")
    ])
```

View the model summary

```
In [ ]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 12)	108
dense_1 (Dense)	(None, 1)	13
Total params: 121 (484.00 Byte)		
Trainable params: 121 (484.00 Byte)		
Non-trainable params: 0 (0.00 Byte)		

## Train the model

- Compile the model with optimizer, loss function and metrics
- Use the fit function to return the run history.

```
In [ ]: model.compile(SGD(lr = .003), "binary_crossentropy", metrics=["accuracy"])
run_hist_1 = model.fit(X_train_norm, y_train, validation_data=(X_test_norm, y_t
```

WARNING:abs1:`lr` is deprecated in Keras optimizer, please use `learning\_rate` or use the legacy optimizer, e.g.,tf.keras.optimizers.legacy.SGD.

Epoch 1/200

```
18/18 [=====] - 1s 14ms/step - loss: 0.6772 - accu
racy: 0.6441 - val_loss: 0.7028 - val_accuracy: 0.6042
```

Epoch 2/200

```
18/18 [=====] - 0s 5ms/step - loss: 0.6607 - accur
acy: 0.6684 - val_loss: 0.6846 - val_accuracy: 0.6302
```

Epoch 3/200

```
18/18 [=====] - 0s 5ms/step - loss: 0.6457 - accur
acy: 0.6753 - val_loss: 0.6685 - val_accuracy: 0.6562
```

Epoch 4/200

```
18/18 [=====] - 0s 4ms/step - loss: 0.6319 - accur
acy: 0.6858 - val_loss: 0.6539 - val_accuracy: 0.6562
```

Epoch 5/200

```
18/18 [=====] - 0s 4ms/step - loss: 0.6196 - accur
acy: 0.6927 - val_loss: 0.6407 - val_accuracy: 0.6615
```

Epoch 6/200

12/12 [-----] - Ac Amc/ctdn - Incc. 0 6022 - accuip

```
In [ ]: ## Like we did for the Random Forest, we generate two kinds of predictions
        # One is a hard decision, the other is a probabilistic score.
```

```
y_pred_class_nn_1 = (model.predict(X_test_norm) > 0.5).astype("int32")
y_pred_prob_nn_1 = model.predict(X_test_norm)
```

6/6 [=====] - 0s 3ms/step

6/6 [=====] - 0s 2ms/step

```
In [ ]: # Let's check out the outputs to get a feel for how keras apis work.  
        y_pred_class_nn_1[:10]
```

[illegible]

```
In [ ]: y_pred_prob_nn_1[:10]
```

```
Out[16]: array([[0.6247541 ],
                [0.78174627],
                [0.2992891 ],
                [0.30132923],
                [0.17397451],
                [0.46647456],
                [0.02336803],
                [0.3432233 ],
                [0.92992723],
                [0.11837403]], dtype=float32)
```


Create the plot\_roc function

```
In [ ]: def plot_roc(y_test, y_pred, model_name):
        fpr, tpr, thr = roc_curve(y_test, y_pred)
        fig, ax = plt.subplots(figsize=(8, 8))
        ax.plot(fpr, tpr, 'k-')
        ax.plot([0, 1], [0, 1], 'k--', linewidth=.5) # roc curve for random model
        ax.grid(True)
        ax.set(title='ROC Curve for {} on PIMA diabetes problem'.format(model_name),
              xlim=[-0.01, 1.01], ylim=[-0.01, 1.01])
```

Evaluate the model performance and plot the ROC CURVE

```
In [ ]: print('accuracy is {:.3f}'.format(accuracy_score(y_test,y_pred_class_nn_1)))  
print('roc-auc is {:.3f}'.format(roc_auc_score(y_test,y_pred_prob_nn_1)))  
  
plot_roc(y_test, y_pred_prob_nn_1, 'NN')
```

accuracy is 0.745  
roc-auc is 0.817




Plot the training loss and the validation loss over the different epochs and see how it looks

```
In [ ]: run_hist_1.history.keys()
```

```
Out[19]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
In [ ]: fig, ax = plt.subplots()
ax.plot(run_hist_1.history["loss"], 'r', marker='.', label="Train Loss")
ax.plot(run_hist_1.history["val_loss"], 'b', marker='.', label="Validation Loss")
ax.legend()
```

Out[20]: <matplotlib.legend.Legend at 0x7f4fd6f10250>



What is your interpretation about the result of the train and validation loss?

***Based on the graph above, it can be observed that the loss for validation is higher compared to the training. To start, training loss is an indication how the model is handling or fitting the training data; wherein validation loss indicates how the model is fitting the new data.***

***Aside from that, it can be interpreted that the model is overfitting, since the model performs well in training data but performs rather poorly in testing data. At some the 25th epoch, it did not show any signs of underfitting or overfitting. However, as the model was trained for a long period; hence, why it overfitted.***

### Supplementary Activity

- Build a model with two hidden layers, each with 6 nodes
- Use the "relu" activation function for the hidden layers, and "sigmoid" for the final layer
- Use a learning rate of .003 and train for 1500 epochs
- Graph the trajectory of the loss functions, accuracy on both train and test set
- Plot the roc curve for the predictions
- Use different learning rates, numbers of epochs, and network structures.

- Plot the results of training and validation loss using different learning rates, number of epochs and network structures
- Interpret your result

```
In [ ]: filepath2 = "/content/drive/MyDrive/Datasci3/breast-cancer.csv"
breast_df = pd.read_csv(filepath2)
```

```
In [ ]: print(breast_df.shape)
breast_df.sample(5)
```

(569, 32)

Out[22]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_m
<b>491</b>	91376702	B	17.85	13.23	114.60	992.1	0.0
<b>233</b>	88206102	M	20.51	27.81	134.40	1319.0	0.0
<b>323</b>	895100	M	20.34	21.51	135.90	1264.0	0.1
<b>330</b>	896839	M	16.03	15.51	105.80	793.2	0.0
<b>506</b>	91544001	B	12.22	20.04	79.47	453.1	0.1

5 rows × 32 columns





```
In [ ]: breast_df.dtypes
```

```
Out[23]: id                int64
diagnosis                object
radius_mean             float64
texture_mean            float64
perimeter_mean          float64
area_mean               float64
smoothness_mean         float64
compactness_mean        float64
concavity_mean          float64
concave points_mean     float64
symmetry_mean           float64
fractal_dimension_mean  float64
radius_se               float64
texture_se              float64
perimeter_se            float64
area_se                 float64
smoothness_se           float64
compactness_se          float64
concavity_se            float64
concave points_se       float64
symmetry_se             float64
fractal_dimension_se    float64
radius_worst            float64
texture_worst           float64
perimeter_worst         float64
area_worst              float64
smoothness_worst        float64
compactness_worst       float64
concavity_worst         float64
concave points_worst    float64
symmetry_worst          float64
fractal_dimension_worst float64
dtype: object
```

```
In [ ]: breast_df = breast_df.drop('id', axis=1)
```

```
In [ ]: breast_df['diagnosis'].astype(str)
```

```
Out[25]: 0      M
1      M
2      M
3      M
4      M
..
564    M
565    M
566    M
567    M
568    B
Name: diagnosis, Length: 569, dtype: object
```

```
In [ ]: breast_df['diagnosis'].value_counts()
```

```
Out[26]: B    357  
        M    212  
        Name: diagnosis, dtype: int64
```

```
In [ ]: from sklearn.preprocessing import LabelEncoder  
        le = LabelEncoder()  
  
        breast_df['diagnosis'] = le.fit_transform(breast_df['diagnosis'])
```

```
In [ ]: breast_df['diagnosis'].value_counts()
```

```
Out[28]: 0    357  
        1    212  
        Name: diagnosis, dtype: int64
```

```
In [ ]: breast_df['diagnosis'].astype(int)
```

```
Out[29]: 0      1  
        1      1  
        2      1  
        3      1  
        4      1  
        ..  
        564    1  
        565    1  
        566    1  
        567    1  
        568    0  
        Name: diagnosis, Length: 569, dtype: int64
```

```
In [ ]: X2 = breast_df.drop('diagnosis', axis=1)  
        y2 = breast_df["diagnosis"]
```

```
In [ ]: X2_train, X2_test, y2_train, y2_test = train_test_split(X2, y2, test_size=0.25,
```

```
In [ ]: normalizer = StandardScaler()  
        X_train_norm2 = normalizer.fit_transform(X2_train)  
        X_test_norm2 = normalizer.transform(X2_test)
```

- Build a model with two hidden layers, each with 6 nodes
- Use the "relu" activation function for the hidden layers, and "sigmoid" for the final layer

```
In [ ]: model2 = Sequential([  
        Dense(6, input_shape=(30,), activation="relu"),  
        Dense(6, input_shape=(30,), activation="relu"),  
        Dense(1, activation="sigmoid")  
    ])
```

```
In [ ]: model2.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 6)	186
dense_3 (Dense)	(None, 6)	42
dense_4 (Dense)	(None, 1)	7

=====  
Total params: 235 (940.00 Byte)  
Trainable params: 235 (940.00 Byte)  
Non-trainable params: 0 (0.00 Byte)  
=====

Use a learning rate of .003 and train for 1500 epochs

```
In [ ]: model2.compile(SGD(lr = .003), "binary_crossentropy", metrics=["accuracy"])  
run_hist_2 = model2.fit(X_train_norm2, y2_train, validation_data=(X_test_norm2,
```

WARNING:abs1:`lr` is deprecated in Keras optimizer, please use `learning\_rate` or use the legacy optimizer, e.g.,tf.keras.optimizers.legacy.SGD.

Epoch 1/1500

14/14 [=====] - 2s 37ms/step - loss: 0.7157 - accuracy: 0.3568 - val\_loss: 0.7274 - val\_accuracy: 0.3566

Epoch 2/1500

14/14 [=====] - 0s 10ms/step - loss: 0.7032 - accuracy: 0.3451 - val\_loss: 0.7113 - val\_accuracy: 0.3986

Epoch 3/1500

14/14 [=====] - 0s 7ms/step - loss: 0.6914 - accuracy: 0.3826 - val\_loss: 0.6963 - val\_accuracy: 0.4266

Epoch 4/1500

14/14 [=====] - 0s 10ms/step - loss: 0.6804 - accuracy: 0.4390 - val\_loss: 0.6813 - val\_accuracy: 0.4755

Epoch 5/1500

14/14 [=====] - 0s 11ms/step - loss: 0.6694 - accuracy: 0.4695 - val\_loss: 0.6658 - val\_accuracy: 0.5105

Epoch 6/1500

14/14 [=====] - 0s 8ms/step - loss: 0.6579 - accuracy: 0.5105

```
In [ ]: y_pred_class_nn_2 = (model2.predict(X_test_norm2) > 0.5).astype("int32")  
y_pred_prob_nn_2 = model2.predict(X_test_norm2)
```


5/5 [=====] - 0s 3ms/step

5/5 [=====] - 0s 2ms/step

```
In [ ]: print('accuracy is {:.3f}'.format(accuracy_score(y2_test,y_pred_class_nn_2)))
print('roc-auc is {:.3f}'.format(roc_auc_score(y2_test,y_pred_prob_nn_2)))

plot_roc(y2_test, y_pred_prob_nn_2, 'NN')
```

accuracy is 0.972  
roc-auc is 0.986




*It can be seen that this graph can accurately predict values with a high accuracy. Considering that its accuracy is 97.2, and it's roc-auc is 98.6, it can be inferred that this model is performing fairly since it does not experience high difficulty in distinguishing both classes. However, it cannot be overlooked that this model is still sensitive to small changes (new data).*

```
In [ ]: run_hist_2.history.keys()
```

```
Out[41]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
In [ ]: fig, ax = plt.subplots()
ax.plot(run_hist_2.history["loss"], 'r', marker='.', label="Train Loss")
ax.plot(run_hist_2.history["val_loss"], 'b', marker='.', label="Validation Loss")
ax.legend()
```

Out[42]: <matplotlib.legend.Legend at 0x7f4fd6d7d930>



***Based on the given required parameters, such parameters is not a fit for this model. It is highly observable how this model overfitted as the training period is prolonged. Therefore, it can be concluded that this model, with the given parameters, is not suitable in distinguishing between the two classes.***

- Use different learning rates, numbers of epochs, and network structures.

```
In [ ]: model3 = Sequential([
    Dense(10, input_shape=(30,), activation="relu"),
    Dense(10, input_shape=(30,), activation="relu"),
    Dense(10, input_shape=(30,), activation="relu"),
    Dense(10, input_shape=(30,), activation="relu"),
    Dense(1, activation="sigmoid")
])
```

```
In [ ]: model3.compile(SGD(lr = .001), "binary_crossentropy", metrics=["accuracy"])
run_hist_3 = model3.fit(X_train_norm2, y2_train, validation_data=(X_test_norm2,
```

WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning\_rate` or use the legacy optimizer, e.g.,tf.keras.optimizers.legacy.SGD.

Epoch 1/500

14/14 [=====] - 2s 26ms/step - loss: 0.7206 - accuracy: 0.6127 - val\_loss: 0.6999 - val\_accuracy: 0.6923

Epoch 2/500

14/14 [=====] - 0s 7ms/step - loss: 0.7020 - accuracy: 0.6268 - val\_loss: 0.6824 - val\_accuracy: 0.6993

Epoch 3/500

14/14 [=====] - 0s 6ms/step - loss: 0.6862 - accuracy: 0.6362 - val\_loss: 0.6674 - val\_accuracy: 0.7203

Epoch 4/500

14/14 [=====] - 0s 7ms/step - loss: 0.6726 - accuracy: 0.6901 - val\_loss: 0.6538 - val\_accuracy: 0.7762

Epoch 5/500

14/14 [=====] - 0s 8ms/step - loss: 0.6606 - accuracy: 0.7465 - val\_loss: 0.6410 - val\_accuracy: 0.8252

Epoch 6/500

14/14 [=====] - 0s 7ms/step - loss: 0.6493 - accuracy: 0.7762 - val\_loss: 0.6299 - val\_accuracy: 0.8523

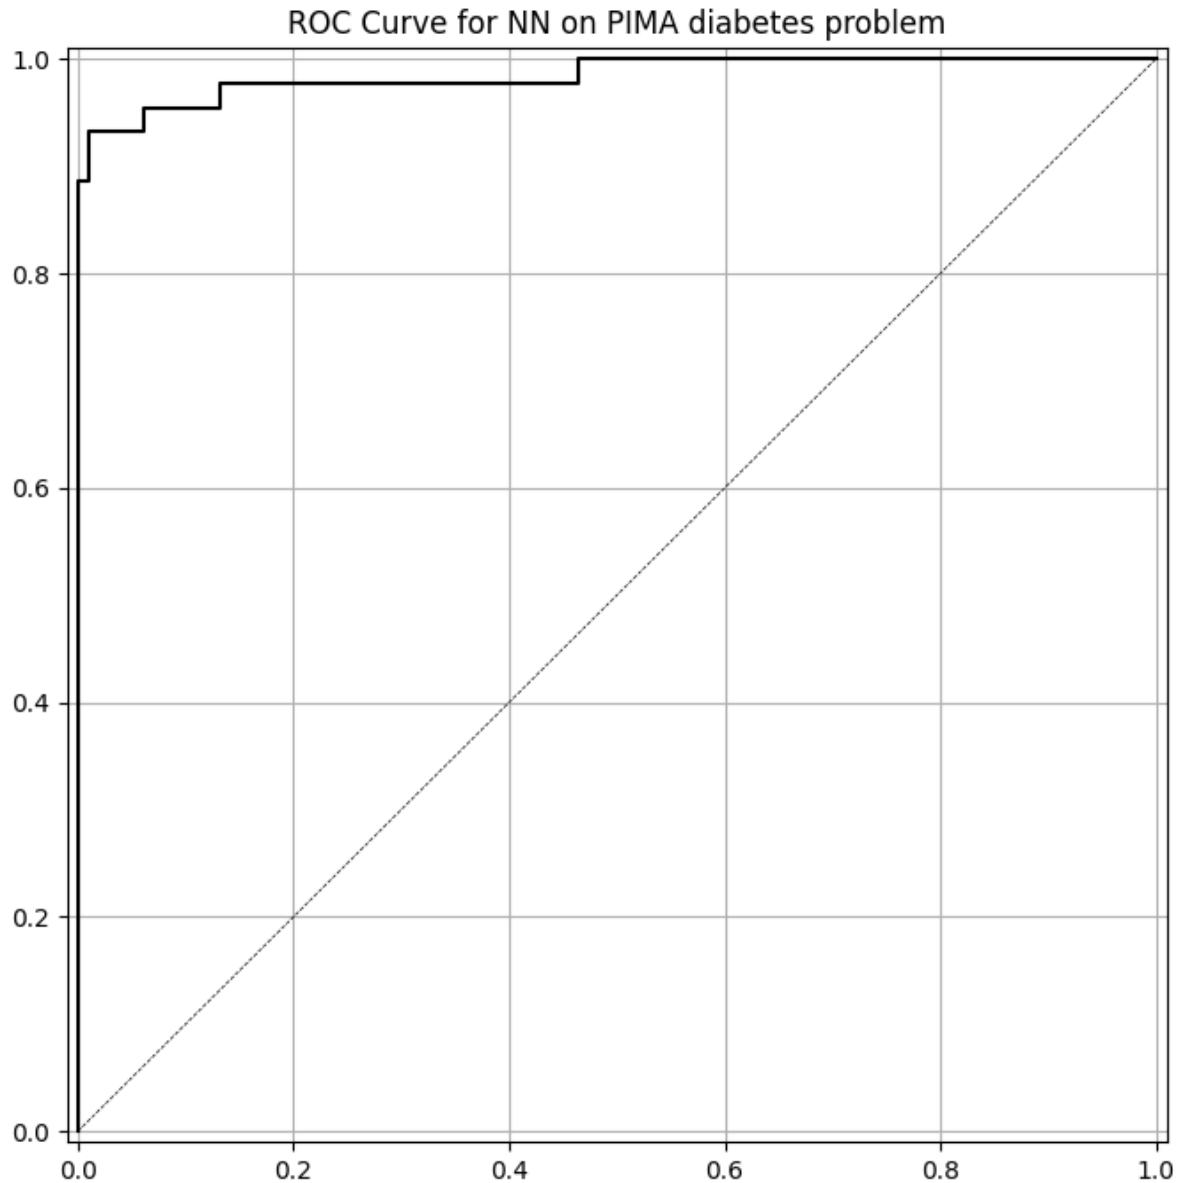
```
In [ ]: y_pred_class_nn_3 = (model3.predict(X_test_norm2) > 0.5).astype("int32")
y_pred_prob_nn_3 = model3.predict(X_test_norm2)
```

5/5 [=====] - 0s 7ms/step

5/5 [=====] - 0s 3ms/step

```
In [ ]: print('accuracy is {:.3f}'.format(accuracy_score(y2_test,y_pred_class_nn_3)))  
print('roc-auc is {:.3f}'.format(roc_auc_score(y2_test,y_pred_prob_nn_3)))  
  
plot_roc(y2_test, y_pred_prob_nn_3, 'NN')
```

accuracy is 0.972  
roc-auc is 0.985



***With the given parameters, it can be observed that the model is still quite having the trouble in distinguishing between the two classes. Albeit its accuracy is 97.2, the threshold in reaching 1.0 was close, and being near 1.0 indicates that this is a fairly-well model.***


***The curve did not show any significant difference compared with the first ROC curve.***

```
In [ ]: run_hist_3.history.keys()
```

```
Out[48]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
In [ ]: fig, ax = plt.subplots()
ax.plot(run_hist_3.history["loss"], 'r', marker='.', label="Train Loss")
ax.plot(run_hist_3.history["val_loss"], 'b', marker='.', label="Validation Loss")
ax.legend()
```

```
Out[49]: <matplotlib.legend.Legend at 0x7f4fd1b2c4f0>
```



***It can be seen that the parameters used also causes the model to overfit, after a lot of training periods, the model slowly overfitted. At the range of 120-130th epoch, it showed signs of overfitting. Ergo, the parameters did not fit the model, as it still showed signs of being sensitive to new data.***

```
In [ ]: model4 = Sequential([
    Dense(6, input_shape=(30,), activation="relu"),
    Dense(6, input_shape=(30,), activation="relu"),
    Dense(6, input_shape=(30,), activation="relu"),
    Dense(6, input_shape=(30,), activation="relu"),
    Dense(1, activation="sigmoid")
])
```



```
In [ ]: model4.compile(SGD(lr = .001), "binary_crossentropy", metrics=["accuracy"])
run_hist_4 = model4.fit(X_train_norm2, y2_train, validation_data=(X_test_norm2,
```

WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning\_rate` or use the legacy optimizer, e.g.,tf.keras.optimizers.legacy.SGD.

Epoch 1/500

14/14 [=====] - 3s 48ms/step - loss: 0.7254 - accuracy: 0.4319 - val\_loss: 0.7208 - val\_accuracy: 0.4895

Epoch 2/500

14/14 [=====] - 0s 27ms/step - loss: 0.7097 - accuracy: 0.4930 - val\_loss: 0.7040 - val\_accuracy: 0.5664

Epoch 3/500

14/14 [=====] - 0s 23ms/step - loss: 0.6983 - accuracy: 0.5469 - val\_loss: 0.6917 - val\_accuracy: 0.6084

Epoch 4/500

14/14 [=====] - 0s 10ms/step - loss: 0.6905 - accuracy: 0.6080 - val\_loss: 0.6845 - val\_accuracy: 0.6643

Epoch 5/500

14/14 [=====] - 0s 11ms/step - loss: 0.6860 - accuracy: 0.6549 - val\_loss: 0.6789 - val\_accuracy: 0.7063

Epoch 6/500

14/14 [=====] - 0s 15ms/step - loss: 0.6824 - accuracy: 0.6824 - val\_loss: 0.6789 - val\_accuracy: 0.7063

```
In [ ]: y_pred_class_nn_4 = (model4.predict(X_test_norm2) > 0.5).astype("int32")
y_pred_prob_nn_4 = model4.predict(X_test_norm2)
```

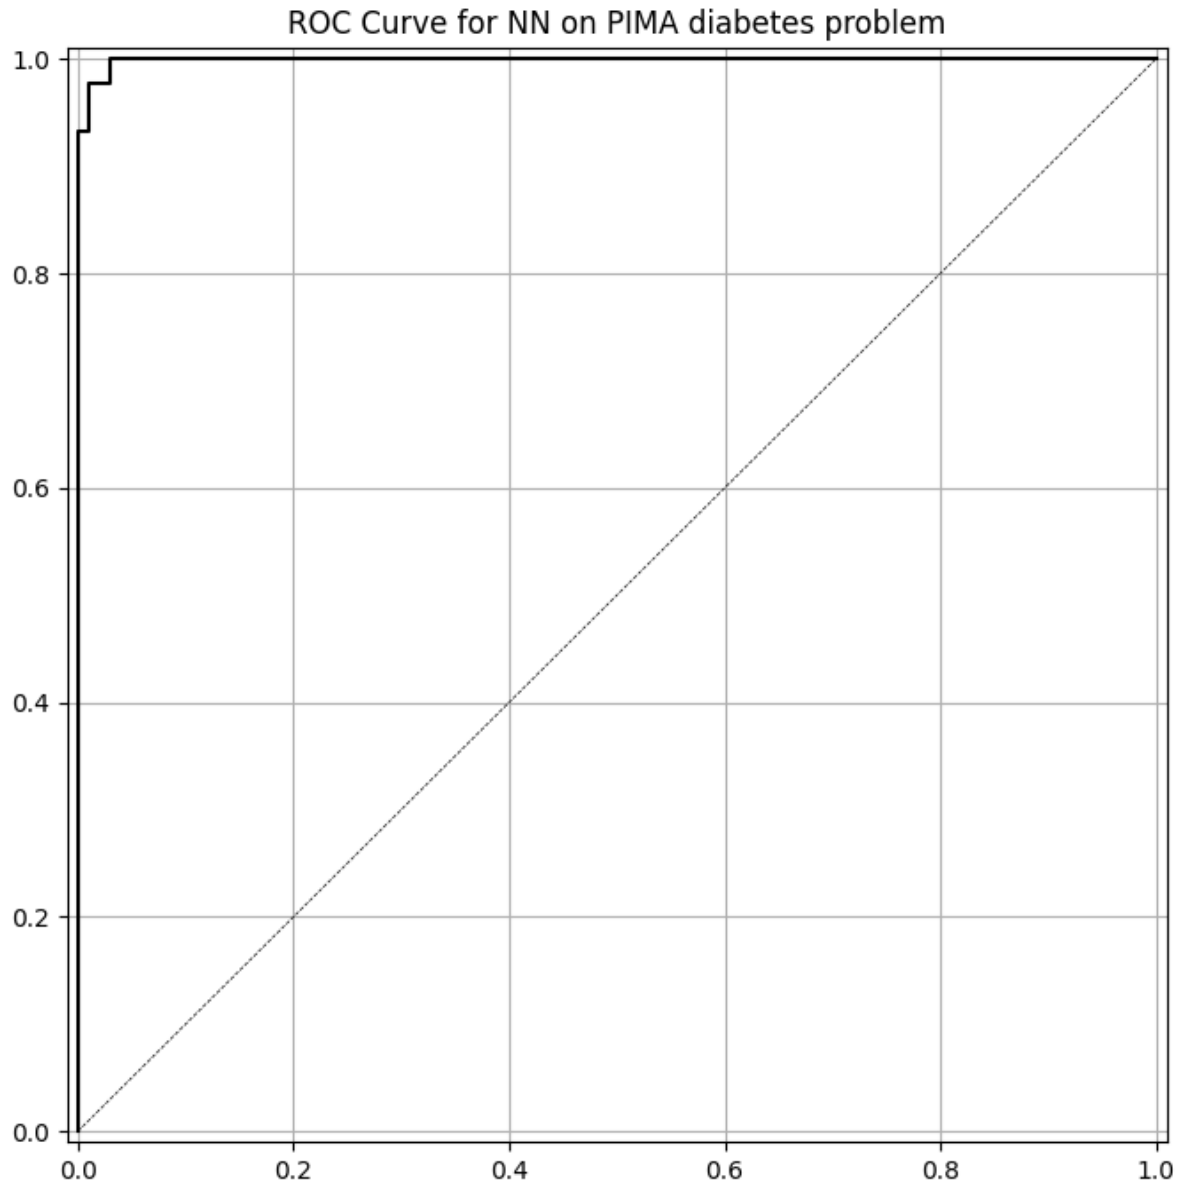
5/5 [=====] - 0s 3ms/step

5/5 [=====] - 0s 3ms/step

```
In [ ]: print('accuracy is {:.3f}'.format(accuracy_score(y2_test,y_pred_class_nn_4)))
print('roc-auc is {:.3f}'.format(roc_auc_score(y2_test,y_pred_prob_nn_4)))

plot_roc(y2_test, y_pred_prob_nn_4, 'NN')
```

accuracy is 0.986  
roc-auc is 0.999



*Here, it can be seen that the error is very minimal, and the roc curve is close to have a perfect trend. Considering that the accuracy is 98.6% and the roc-auc is 99%, this parameters served as a great fit for this model. With 4 hidden layers, 1 final layer, .001 learning rate, and a total of 500 epochs. The same activation function and function of the final layer was used with the former network structure.*


*Hence, this model was able to distinguish between the classes outstandingly since the roc cruve leans towards the corner (as it's an indicator of whether the model is performing well or not.)*

```
In [ ]: run_hist_4.history.keys()
```

```
Out[270]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
In [ ]: fig, ax = plt.subplots()
ax.plot(run_hist_4.history["loss"], 'r', marker='.', label="Train Loss")
ax.plot(run_hist_4.history["val_loss"], 'b', marker='.', label="Validation Loss")
ax.legend()
```

```
Out[271]: <matplotlib.legend.Legend at 0x7f4fc6bd4100>
```



***Based on the figure, it can be stated that 4 hidden layer with relu activations, and 1 final layer with a sigmoid function, and alongwith a .001 learning rate, BreastCancer\_df proved to be effective with this network structure and parameters. It can be seen that it only slightly overfits, but the trend of the plot is clean and promising. This model can perform well given the set of parameters - comparing it to the others, this model can handle new data and is not sensitive to changes, as it does not overfits completely.***

## Conclusion

***It can be concluded that this activity was a similar in training a machine learning model. There are only different parameters and some functions wasn't performed in machine learning. Moreover, I was able to grasp the concept in building and training a neural network (which is not so different from ML) and evaluating and plotting the model using training and validation loss. Formerly, evaluation and plotting of the model was reliant on the classification report. However, in deep learning, training and validation loss was used.***

***Other than that, this activity served as a refresher in training models, and another thing that I noticed is that this activity haven't utilized grid search yet, which was always***